

Microsoft Agent 365 Security PoC Runbook

Setup, tests, and result checks for the six Agent 365 PoC use cases across the Observe, Govern, and Secure pillars.

Setup, tests, and result checks for the six Agent 365 PoC use cases across the Observe, Govern, and Secure pillars.

Least-privilege model used in this runbook

Every chapter has a **Roles** section that splits access into:

- **Setup / configuration role** — the elevated role required only to make a change (create a policy, enable a feature, publish a label). Assign Active for the PoC window; PIM-elevate if needed.
- **Read-only / validation role** — the least-privilege role sufficient to view configuration, evidence, and logs. Use this for reviewers, auditors, and everyone who does not need to change anything.

Anyone verifying results or reviewing evidence should use the read-only role. Setup roles stay with the admin who is actually making the change.

Structure

Pillar	Use case	Folder
—	Chapter 0 – Prerequisites and PoC preparation	chapter-0-prerequisites
Observe	UC1 – Agent Discovery	chapter-uc1-agent-discovery
Govern	UC2 – Identity & Ownership	chapter-uc2-identity-ownership
Govern	UC3 – Least-Privilege Access	chapter-uc3-least-privilege
Secure (Purview)	UC4 – Sensitive Data Protection	chapter-uc4-sensitive-data-protection
Secure (Defender)	UC5 – Threat Detection & Protection	chapter-uc5-threat-detection
Observe / Govern	UC6 – Lifecycle & Audit	chapter-uc6-lifecycle-audit

Combined PDF at the root: `Agent-365-Security-PoC-Runbook.pdf` .

Portals used

Portal	URL	Purpose
Microsoft 365 admin center	admin.cloud.microsoft	Agents, registry, approvals, connected platforms
Microsoft Entra admin center	entra.microsoft.com	Agent ID, Conditional Access, access packages, lifecycle workflows, sign-in logs
Microsoft Purview portal	purview.microsoft.com	Labels, DLP, audit, Communication Compliance, IRM
Microsoft Defender portal	security.microsoft.com	AI agent inventory, alerts, Advanced Hunting
Power Platform admin center	admin.powerplatform.microsoft.com	Copilot Studio environment, security integrations
Copilot Studio	copilotstudio.microsoft.com	Build and publish Copilot Studio agents

Final readiness checklist

Before the joint test session, confirm:

- ☐ PoC agents are published and accessible to the test user.
- ☐ Agent 365 is connected in Defender.
- ☐ Defender Security for AI agents is enabled.
- ☐ Power Platform admin is available if Copilot Studio real-time protection is in scope.
- ☐ Microsoft 365 app connector is connected if near-real-time detections are required.
- ☐ Purview Audit is enabled.
- ☐ Sensitivity labels are published and visible to test users.
- ☐ Confidential and General files are uploaded and labeled.
- ☐ DLP policy for Copilot / Copilot Chat is enabled.
- ☐ DLP policy for Exchange email is enabled.
- ☐ Test user has SharePoint access and can invoke the PoC agent.
- ☐ Propagation windows have passed.
- ☐ Defender and Purview admins are available during the test window.

Chapter 0 - Prerequisites and PoC preparation

Baseline steps and readiness checks before running any of the six use cases.

What this chapter covers

- Which use cases are in scope for this PoC.
- Access and licensing.
- Test identities (least-privilege split).
- Which portals every team needs.
- Test agents to prepare.
- SharePoint content for Purview tests (UC4).
- Propagation windows.

Least-privilege model

Every chapter in this runbook uses the same pattern:

- **Setup / configuration role** - elevated role required only to make a change. Assign Active for the PoC window.
- **Read-only / validation role** - least-privilege role sufficient to view configuration and evidence. Use this for reviewers, auditors, and anyone who does not need to change anything.

Grant read-only roles first. Only grant setup roles to whoever will actually make the change.

0.1 Confirm PoC scope

Agree upfront which use cases are in scope for this PoC:

Use case	Pillar	Included in this guide
UC1 - Agent Discovery	Observe	Yes
UC2 - Identity & Ownership	Govern	Yes
UC3 - Least-Privilege Access	Govern	Yes
UC4 - Sensitive Data Protection (Purview)	Secure	Yes
UC5 - Threat Detection & Protection (Defender)	Secure	Yes
UC6 - Lifecycle & Audit	Observe / Govern	Yes

0.2 Confirm access and licensing

Confirm the PoC team has:

- Agent 365 trial license or equivalent Agent 365 access for the PoC.
- Access to the Microsoft 365 admin center (<<https://admin.cloud.microsoft>>).

- Access to the Microsoft Entra admin center (<<https://entra.microsoft.com>>).
- Access to the Microsoft Purview portal (<<https://purview.microsoft.com>>).
- Access to the Microsoft Defender portal (<<https://security.microsoft.com>>).
- Access to the Power Platform admin center (if Copilot Studio is in scope).

0.3 Roles at a glance

Assign these Active for the PoC window. Full role detail is in each chapter.

Use case	Setup role	Read-only / validation role
UC1 - Agent Discovery	AI Administrator	AI Reader or Global Reader
UC2 - Identity & Ownership	Agent ID Administrator or Cloud Application Administrator (manage), Agent ID Developer (blueprints), Attribute Definition Administrator, Attribute Assignment Administrator	No admin role needed to view (any Microsoft Entra user)
UC3 - Least-Privilege Access	Conditional Access Administrator, Identity Governance Administrator, Lifecycle Workflows Administrator	Global Reader, Reports Reader, Security Reader
UC4 - Sensitive Data Protection	Compliance Administrator, Information Protection Administrator, Purview Workload Content Admin, Communication Compliance Administrator, Insider Risk Management Administrator	Audit Reader, Global Reader, Compliance Reader
UC5 - Threat Detection & Protection	Security Administrator, Security Operator, Power Platform Administrator (for Copilot Studio integration), Application Administrator or Cloud Application Administrator	Security Reader
UC6 - Lifecycle & Audit	AI Administrator, Compliance Administrator	Audit Reader, AI Reader, Security Reader, Global Reader
Setup only	Global Administrator (initial setup / admin consent only, via PIM / just-in-time)	-

Two rules that save PoC time:

1. Roles assigned as **eligible-only** show up as "you lack permission" to the admin at the moment they try to act. Always assign the role **Active** for the PoC window. Use PIM only for Global Administrator.
2. Split the role assignments explicitly - do not give the admin who executes a change the read-only reviewer's role, and do not give the reviewer the admin role.

0.4 Confirm test identities

Prepare at least two accounts:

Account	Purpose	Roles assigned
Test admin	Configures Agent 365, Entra, Purview, Defender, SharePoint, labels, policies, and agent publishing.	Setup roles per the table in 0.3, scoped to the use cases in scope.
Standard test user	Runs the agent tests. Permission trimming, labels, and DLP evaluate against this user.	Standard user with access to the PoC agent, SharePoint site, and Teams / M365 Copilot as required.
Reviewer / auditor (optional)	Views evidence without making changes.	Read-only / validation roles per the table in 0.3.

0.5 Prepare agents

Before starting UC1:

1. Ensure the PoC agent or agents are published and visible to the intended test users.
2. If using Copilot Studio, publish the agent to Microsoft 365 Copilot and Teams.
3. If using Agent 365 registration, confirm the agent appears in the Agent 365 registry.
4. Exercise the agent with a few baseline prompts so activity exists for logs, hunting, and audit.

0.6 Prepare SharePoint content for Purview tests (UC4)

Create a dedicated SharePoint site or folder for PoC content. A dedicated site is preferred for clean permissions and easy teardown.

Recommended content categories:

Content category	Label	Purpose
Sensitive customer / financial documents	Confidential	Used to prove Copilot grounding is restricted.
Public or broadly shareable operational documents	General	Used as the positive control for allowed grounding.
Financial identifiers in invoice or ledger content	Confidential	Used to prove Exchange DLP blocks outbound email exfiltration.

Sample files are provided for UC4 testing. See [UC4 - Sensitive Data Protection](#) for the file list, labeling guidance, and the samples folder (`chapter-uc4-sensitive-data-protection/samples/`).

All sample files use publicly documented test values (Visa/MC/Amex/Discover sandbox card numbers and Federal Reserve public test routing numbers). No real PII.

0.7 Important propagation windows

Plan for propagation time:

Configuration	Expected delay
Sensitivity label publishing	Up to 24 hours
DLP policy activation	Up to 24 hours
Communication Compliance policy	Up to 24 hours
Insider Risk Management agent policy	Up to 24 hours
Advanced Hunting / audit ingestion	Often 15-30 minutes, sometimes longer
Agent Registry discovery	Up to 24 hours after agent enablement

Chapter UC1 - Agent Discovery

Pillar: Observe

What it proves: you cannot govern what you cannot see. Every agent in the tenant surfaces in one registry, including the ones nobody registered, each with a named, accountable owner.

Roles - least privilege

Task	Role	Notes
Setup / configuration	AI Administrator	Required to enable Agent 365, configure Registry sync for connected platforms, approve agents, and change registry settings. Assign Active for the PoC window.
Read-only / validation	AI Reader (or Global Reader)	Sufficient to view the Agent Registry, list agents, see owners, and export the registry. Use this role for reviewers, auditors, and anyone verifying results.
Optional (SDK agents)	Agent ID Developer	Required for developers onboarding custom-built agents via the Agent 365 SDK. Not required for viewing.

Grant the read-only role first; only grant the setup role to whoever will make configuration changes.

Scope - agent types and discovery paths

The Agent Registry discovers three categories of agent. Each has a different discovery path and surfaces different metadata. Sample agents of each type are provided in `agent-samples/`.

Category	Example platforms	Discovery path	Sample
Microsoft-native	Microsoft 365 Copilot / Agent Builder declarative agents, Copilot Studio, Azure AI Foundry, SharePoint agents	Automatic once Agent 365 is enabled. Agent appears in the Registry after publish (Copilot Studio / Agent Builder) or after agent creation (Foundry). Full propagation up to 24 hours.	<code>agent-samples/copilot-studio/</code> , <code>agent-samples/foundry/</code>
Connected platforms	Databricks Genie, Google Vertex AI, Snowflake Cortex, AWS Bedrock, Salesforce Agentforce	Registry sync configured per platform in <code>admin.cloud.microsoft > Agents > Connected platforms</code> . The source platform's agent inventory is pulled into the Agent Registry on a schedule. See the optional connected-platforms section at the end of this chapter.	-
Custom (Agent 365 SDK)	Customer-built agents in Node.js, Python, or .NET	Onboarded through <code>microsoft/agent365-skills</code> by running <code>a365 setup all</code> . Appears in the Registry once the Blueprint and Entra Agent Identity are provisioned.	<code>agent-samples/sdk/</code>

Each sample folder has a README that walks through the discovery path step by step.

Portal

Microsoft 365 admin center - <<https://admin.cloud.microsoft>> > **Agents**.

Documentation

Topic	Documentation
Agent 365 overview	Overview of Microsoft Agent 365
Agent Registry	Agent Registry in Microsoft Agent 365
Connected platforms	Connected platforms for Agent 365

Prerequisites

Complete [Chapter 0 - Prerequisites](#). In addition:

1. Agent 365 is enabled for the tenant.
2. Frontier features are opted in if Frontier-only capabilities are in scope.
3. At least one Microsoft-native agent exists (Copilot Studio, Foundry, or Microsoft 365 Copilot Agent Builder) so there is content to discover.
4. If connected platforms are in scope (Databricks, Vertex, Snowflake, Bedrock, Agentforce), those tenant-side integrations are ready.

Setup steps

Performed by the AI Administrator.

1. Open <<https://admin.cloud.microsoft>> and go to **Agents**.
2. Confirm **Agent Registry** is enabled.
3. For each connected platform in scope, configure Registry sync per that platform's onboarding page.
4. For custom SDK agents, confirm they have been onboarded via `a365 setup all` from the `microsoft/agent365-skills` toolchain so they appear in the registry.
5. Exercise each agent at least once (send a prompt or invoke a tool) so it registers activity and appears with the expected metadata.

Check result (Setup role)

- Registry contains every in-scope agent.
- Each agent shows a named business owner.
- Connected platform syncs report healthy.

Test - registry inventory

Performed by anyone with **AI Reader**.

1. Open <<https://admin.cloud.microsoft>> > **Agents** > **All agents**.
2. Filter by platform and confirm every in-scope agent is listed.

3. Confirm the owner column is populated for every agent.
4. Export the registry (CSV or JSON) as evidence.

Expected result

- All in-scope agents are visible in the registry within the agreed discovery window (default: 24 hours after enablement).
- Every agent has a business owner.
- No unknown or orphaned agents remain.

Test - Agent Map

Performed by anyone with **AI Reader**.

1. Open the **Agent Map** view in the Agents section.
2. Confirm the map renders agent-to-tool and agent-to-agent interactions.
3. Capture a screenshot as evidence.

Expected result

- Agent Map renders known interactions.
- Discovered agents that were not previously known are triaged with their owner.

Evidence to capture

- Registry export (CSV/JSON) with owner column populated.
- Agent Map screenshot.
- List of any newly discovered / previously unknown agents, with owner assignment status.

Common issues

Symptom	Likely cause	Fix
Agent missing from registry	Agent 365 not enabled for that platform, or agent never invoked	Enable the platform integration; run at least one turn on the agent.
Owner column blank	Sponsor / owner not assigned during agent creation	Assign an owner in the agent's platform (Copilot Studio, Foundry, or Entra Enterprise Applications). Covered in UC2 .
Connected platform not syncing	Registry sync not configured or credentials expired	Reconfigure Registry sync per platform documentation.
SDK agent missing	<code>a365 setup all</code> not completed for that project	Rerun <code>a365 setup all</code> and confirm Blueprint is registered.

Optional - connected platforms (AWS Bedrock, Google Vertex AI, Databricks Genie)

Agent 365 can pull agents from non-Microsoft platforms into the same Registry so they are visible, assignable to owners, and in scope for the governance and audit workstreams. Setup is a one-time per-platform configuration in the Microsoft

365 admin center. All three paths below share the same pattern - prepare the source platform, register a connection, and let Registry sync populate the inventory on its schedule.

Common portal entry point

Microsoft 365 admin center - <<https://admin.cloud.microsoft>> > **Agents** > **Connected platforms** > + **Add platform**.

AWS Bedrock

Step	What to do
1. AWS prerequisites	At least one Bedrock agent exists in an AWS account and region in scope. Create an IAM role or user with least-privilege read access to Bedrock agents (for example <code>bedrock:ListAgents</code> , <code>bedrock:GetAgent</code>).
2. Credentials	Prepare the credentials or federation that Agent 365 will use to call Bedrock (IAM role ARN with trust policy, or access key for the IAM user).
3. In Agent 365	Open Connected platforms > + Add platform > AWS Bedrock . Enter the AWS account ID, region(s), and credentials or federation details. Save.
4. Validate	Confirm connector status is Connected . Wait for the first sync (typically within an hour).
5. Verify in Registry	Bedrock agents appear under Agents > All agents with platform = AWS Bedrock, the AWS agent ID, and a configurable business owner.

Documentation:

- [Amazon Bedrock Agents](#)
- [IAM for Bedrock](#)
- [Agent 365 - Connected platforms](#)

Google Vertex AI

Step	What to do
1. Google Cloud prerequisites	At least one Vertex AI agent exists in a Google Cloud project in scope. Enable the Vertex AI API on the project.
2. Service account	Create a Google Cloud service account with least-privilege read access to Vertex AI (for example <code>aiplatform.agents.list</code> , <code>aiplatform.agents.get</code>). Download the service-account key or configure Workload Identity Federation with Entra.
3. In Agent 365	Open Connected platforms > + Add platform > Google Vertex AI . Enter the GCP project ID, region, and the service-account credentials or federation. Save.
4. Validate	Confirm connector status is Connected . Wait for the first sync.
5. Verify in Registry	Vertex agents appear with platform = Google Vertex AI, the Vertex agent/project ID, and an owner you can edit.

Documentation:

- [Vertex AI Agents overview](#)
- [Service accounts in Google Cloud](#)
- [Workload Identity Federation with Microsoft Entra](#)
- [Agent 365 - Connected platforms](#)

Databricks Genie

Step	What to do
1. Databricks prerequisites	At least one Genie space / agent exists in the Databricks workspace(s) in scope. Confirm the workspace URL and that the Genie product is enabled.
2. Databricks credentials	Create a Databricks service principal (or personal access token for a bootstrap) with least-privilege read scope on Genie resources.
3. In Agent 365	Open Connected platforms > + Add platform > Databricks Genie . Enter the Databricks workspace URL and credentials. Save.
4. Validate	Confirm connector status is Connected . Wait for the first sync (typically within an hour).
5. Verify in Registry	Genie agents appear with platform = Databricks Genie, the Databricks workspace, and owners you can assign.

Documentation:

- [Databricks AI/BI Genie overview](#)
- [Databricks service principals](#)
- [Agent 365 - Connected platforms](#)

Test - connected agents in the Registry

Performed by anyone with **AI Reader**.

1. Open <<https://admin.cloud.microsoft>> > **Agents** > **All agents**.
2. Filter by platform (AWS Bedrock, Google Vertex AI, Databricks Genie).
3. Confirm every expected connected agent is listed.
4. Confirm each has an assignable owner.
5. Export the filtered list as evidence.

Common issues

Symptom	Likely cause	Fix
Connector status = Error	Invalid credentials or missing permission on the source platform	Verify the IAM / service-account / service-principal permissions match the connector's requirements.
No agents sync	Connector connected but source platform has no agents, or Registry sync has not run yet	Confirm the source has agents; wait for the next sync cycle.
Partial sync	Permission scoped too narrowly on the source platform	Grant list + get permissions across the agent resources.
Agent appears without owner	Owner is set in Agent 365, not inherited from the source	Assign a business owner in the Agent 365 Registry UI.

Optional - programmatic discovery with Microsoft Graph

The UI in the Microsoft 365 admin center is the fastest path for a visual review, but Microsoft Graph exposes the same inventory programmatically. Use this when you need scheduled exports, external reporting, CMDB sync, or automated reconciliation against a source of truth.

What you can retrieve

Signal	Where to get it in Graph
List of agents in the tenant (Copilot Studio / M365 Copilot declarative, Foundry, SharePoint, SDK) with display name, owner, publishing state	<code>/copilot/admin/agents</code> and related Copilot admin endpoints (beta)
Agent identity in the directory (Entra Agent ID service principal)	<code>/servicePrincipals</code> filtered to the agent subtype
Owners and sponsors of an agent	<code>/servicePrincipals/{id}/owners</code> , <code>/applications/{id}/owners</code>
Custom security attributes on an agent identity	<code>/servicePrincipals/{id}? \$select=customSecurityAttributes</code>
Conditional Access policies that target agent identities	<code>/identity/conditionalAccess/policies</code>
Sign-in activity for an agent (allow/deny, including CA decisions)	<code>/auditLogs/signIns</code> filtered by <code>appId</code> , service-principal filters
Audit of lifecycle actions on an agent (create, update, block, delete)	<code>/auditLogs/directoryAudits</code> and Microsoft Purview Audit APIs
Copilot interaction records (for UC4/UC6 cross-check)	Microsoft Purview Audit via Microsoft 365 Management Activity API, or Graph audit APIs where available

Minimum required Graph permissions (least-privilege)

Scope	Why
Application.Read.All (delegated or app)	Read agent applications and service principals.
Directory.Read.All (delegated)	Read owners, sponsors, and directory metadata.
CustomSecAttributeAssignment.Read.All (delegated)	Read custom security attribute values on agent identities (requires Attribute Assignment Reader).
Policy.Read.All (delegated)	Read Conditional Access policies that target agent identities.
AuditLog.Read.All (delegated or app)	Read directory audit and sign-in logs.
AgentApplication.Read.All / Copilot admin scopes (preview)	Read Copilot/agent inventory via the Copilot admin endpoints. Scope names and surface are preview; confirm against current docs before relying on them in production.

Prefer delegated permissions scoped to a read-only account. Only grant application permissions for an unattended service that needs to run outside a user context.

Example - list agent service principals

```
GET https://graph.microsoft.com/v1.0/servicePrincipals?$filter=servicePrincipalType eq 'Application'&$select=id,displayName,appId,servicePrincipalType,tags
Authorization: Bearer <token>
```

To focus on agent identities only, filter further by display name convention (for example names ending in - AgentIdentity) or by a known custom security attribute set:

```
GET https://graph.microsoft.com/beta/servicePrincipals?
$count=true&$filter=customSecurityAttributes/AgentGovernance/Project eq 'Agent365PoC'&$select=id,displayName,appId,owners
ConsistencyLevel: eventual
Authorization: Bearer <token>
```

Example - get sign-in activity for one agent

```
GET https://graph.microsoft.com/v1.0/auditLogs/signIns?$filter=appId eq '<agent-app-id>'&$top=50
Authorization: Bearer <token>
```

Returns the service-principal sign-ins for the agent, including Conditional Access decisions. Used in [UC3](#) and [UC6](#).

Example - read custom security attributes on an agent

```
GET https://graph.microsoft.com/v1.0/servicePrincipals/{id}?select=id,displayName,customSecurityAttributes
Authorization: Bearer <token>
```

Documentation

Topic	Documentation
Microsoft Graph overview	Microsoft Graph REST API overview
Service principals in Graph	servicePrincipal resource type
List service principals	List servicePrincipals
Custom security attributes via Graph	Manage custom security attributes using Microsoft Graph
Audit logs	auditLogRoot: directoryAudits
Sign-in logs	auditLogRoot: signIns
Conditional Access policies	conditionalAccessPolicy resource type
Microsoft 365 Copilot admin APIs	Microsoft 365 Copilot developer documentation and Copilot API reference (preview)
Entra Agent ID (reference)	Microsoft Entra Agent ID
PowerShell with Graph	Microsoft Graph PowerShell SDK

When to use Graph vs the portal

Need	Use
One-off review during the PoC	Portal (< https://admin.cloud.microsoft >)
Scheduled export / CMDB reconciliation	Graph via PowerShell SDK or a lightweight job
Cross-tenant / multi-environment reporting	Graph with application permissions, scoped read-only
Evidence pack for an auditor	Portal export plus Graph <code>/auditLogs/*</code> extracts

Chapter UC2 - Identity & Ownership

Pillar: Govern

What it proves: anonymous agents become governable identities. Every agent gets a directory identity in Microsoft Entra ID with assigned **owners** and **sponsors** and a lifecycle state, so its access can be attributed and revoked like any other workload identity.

Concepts (short)

Microsoft Entra Agent ID has three related objects:

Object	What it is
Agent identity blueprint	The parent definition from which individual agent identities are created. Created by an Agent ID Developer ; that user becomes the owner of the blueprint and its service principal.
Agent identity	A directory identity of type "agent" that an agent uses to call APIs. One blueprint can create many agent identities.
Agent user account (optional)	Used when the agent needs its own Microsoft 365 identity, mailbox, Teams presence, etc. (AI Teammate pattern).

Each agent identity has:

- **Owners** - who handle technical administration of the agent.
- **Sponsors** - who are accountable for the agent's purpose and lifecycle decisions.

Owners and sponsors are shown as a single **Owners and Sponsors** column in the Entra admin center.

Roles - least privilege

Taken directly from [Manage agent identities in your organization](#).

Task	Required role	Notes
View agent identities (least privilege)	Any Microsoft Entra user account	No admin role needed to view. Reviewers and auditors need nothing extra.
Manage agent identities (edit, disable, assign owners/sponsors)	Agent ID Administrator or Cloud Application Administrator	Owners of a specific agent can manage their own agents without these roles.
Create agent blueprints	Agent ID Developer	The creator becomes owner of the blueprint and its service principal.
Create and define custom security attributes (CSA) sets	Attribute Definition Administrator	Required to define the attribute set and keys.
Assign CSA values to agent identities	Attribute Assignment Administrator	Required to write values on an agent identity. Split from the definition role by design.
View ID Protection risk reports for agents	Security Administrator, Security Operator, or Security Reader	Requires Microsoft Entra ID P2 (preview).
Configure Lifecycle Workflows for agents (sponsor change)	Lifecycle Workflows Administrator	Covered in UC3 .

Least privilege for the PoC validation step is: no admin role at all. Any user can view agent identities. Grant the admin roles only to the person who actually configures.

Portal

Microsoft Entra admin center - <<https://entra.microsoft.com>> > **Entra ID** > **Agents**:

- **Agent identities** - the primary view for this chapter.
- **Agent blueprints** - view and manage the parent blueprints.

Custom security attributes are under **Entra ID** > **Protect & secure** > **Custom security attributes**.

Documentation

Topic	Documentation
Entra Agent ID - management	Manage agent identities in your organization
View and filter agent identities	View and filter agent identities in your tenant
Manage blueprints	View and manage agent identity blueprints in your tenant
Inheritable permissions	Configure inheritable permissions for agent identity blueprints
Conditional Access for agents	Conditional Access for Agents in Microsoft Entra
Custom security attributes	Custom security attributes in Microsoft Entra ID
Migrate Copilot Studio agents to Agent ID	Recreate Copilot Studio agents in Microsoft Entra Agent ID

Prerequisites

Complete [Chapter 0 - Prerequisites](#) and [UC1 - Agent Discovery](#). In addition:

1. Each in-scope agent has a nominated **owner** (technical administrator) and **sponsor** (business accountability).
2. If Copilot Studio agents are in scope, Entra Agent Identity is enabled at the Copilot Studio environment level in the Power Platform admin center.

2.1 Setup - enable Entra Agent Identity at the Copilot Studio environment level

Required only if Copilot Studio agents are in scope. Performed by a **Power Platform Administrator** or environment admin.

1. Open <<https://admin.powerplatform.microsoft.com>>.
2. Go to **Environments** > select the environment used for PoC agents.
3. Enable **Entra Agent Identity** for the environment.

Once enabled, every new Copilot Studio agent created in that environment automatically receives an Entra Agent ID. Legacy agents created before enablement continue using traditional app registrations and can be migrated using the docs link above.

2.2 Setup - confirm and assign owners and sponsors

Performed by **Agent ID Administrator** or **Cloud Application Administrator**. (The agent's existing owner can also do this for their own agent without those roles.)

1. Sign in to <<https://entra.microsoft.com>>.
2. Go to **Entra ID** > **Agents** > **Agent identities**.
3. Select each PoC agent identity.
4. On the agent's management page, open **Owners and sponsors**.
5. Confirm the correct **owner** (technical administrator) and **sponsor** (business accountable) are listed. Add any missing entries.

2.3 (Optional) Setup - custom security attributes for governance

Performed by **Attribute Definition Administrator** (step 1-2) then **Attribute Assignment Administrator** (step 3).

1. In Entra, go to **Entra ID > Protect & secure > Custom security attributes > Add attribute set** and create `AgentGovernance`.

2. Define two attributes inside the set:

Attribute	Type	Allowed values	Why
Project	String (predefined)	Agent365PoC	Groups everything in this PoC; one filter pulls every PoC agent.
Environment	String (predefined)	Pilot , Prod	Prevents PoC agents from being mistaken for production.

3. For each PoC agent identity (**Entra ID > Agents > Agent identities** > select agent > **Custom security attributes**), assign:

Attribute	Value
Project	Agent365PoC
Environment	Pilot

2.4 Test - identity and ownership (no admin role required)

Performed by any Microsoft Entra user.

1. Sign in to <https://entra.microsoft.com>.

2. Browse to **Entra ID > Agents > Agent identities**.

3. Use the search box to find an agent by **name** or **object ID**, or add the **Blueprint App ID** filter.

4. (Optional) Select **Choose columns** and add **Owners and Sponsors**, **Blueprint App ID**, **Status**, and **Created On**.

5. For each in-scope PoC agent, confirm:

- The agent identity exists.
- **Status** is active (not disabled).
- **Owners and Sponsors** lists the expected owner and sponsor.
- **Blueprint App ID** matches the expected blueprint (confirms it was created from the correct parent).

6. (Optional) Select an agent and open its details page to review:

- Granted permissions.
- Sign-in logs for the agent.
- Audit log entries for owner / sponsor changes.

7. Cross-check in the Agent 365 Registry (<https://admin.cloud.microsoft> > **Agents**) that the same agent shows the same owner.

8. Export the agent-identities list as evidence (CSV from the view, or via Microsoft Graph - see [UC1](#)).

Expected result

- Every in-scope agent has an agent identity in Entra.
- Each agent identity has at least one named **owner** and one named **sponsor**.

- Each agent identity is linked to its parent **blueprint** (visible via Blueprint App ID).
- Owner is consistent between Entra and the Agent 365 Registry.
- If custom security attributes are used, filtering by `AgentGovernance/Project = Agent365PoC` returns every PoC agent in one list.

Evidence to capture

- Export of the Agent identities list for the PoC (Name, Object ID, Status, Blueprint App ID, Owners and Sponsors, CSA values).
- Screenshot of one agent's details page showing Owners and Sponsors, granted permissions, and sign-in log entries.
- Screenshot of the Agent 365 Registry with matching owner.

Common issues

Symptom	Likely cause	Fix
Agent not visible in Agents > Agent identities	Agent is a legacy pre-enablement Copilot Studio agent, or the environment has not enabled Entra Agent Identity	Enable at the Copilot Studio environment level, or migrate legacy agents using the docs link above.
Agent identity has no sponsor	Sponsor was never assigned during agent creation	Add a sponsor from the agent's details page.
Agent identity has no owner	Owner was never assigned, or the original owner left and ownership was not transferred	Add an owner from the agent's details page, or use the sponsor-change workflow in UC3 .
Owner in Entra does not match owner in Agent 365 Registry	Sync delay, or owner was changed in only one surface	Re-assign in Entra; wait for sync; confirm in the Registry.
Reviewer cannot see CSA values on an agent	Reviewer is missing the ability to read the attribute set	Grant Attribute Assignment Reader on the attribute set, or an equivalent role, to the reviewer.

Chapter UC3 - Least-Privilege Access

Pillar: Govern

What it proves: a compromised agent reaches only what you scoped it to. Broad standing access is replaced by a bounded, just-in-time, revocable grant, and the deny is provable in the sign-in log.

Roles - least privilege

Task	Role	Notes
Create and enforce Conditional Access policies for agents	Conditional Access Administrator	Required to create, edit, and enable CA policies.
Mark an agent as risky for the test (Graph <code>confirmCompromised</code>)	Security Administrator or Security Operator	Required to call Identity Protection APIs.
Create access packages and JIT policies	Identity Governance Administrator	Required for Entitlement Management.
Configure Lifecycle Workflows (sponsor change)	Lifecycle Workflows Administrator	Required only if sponsor-change workflows are in scope.
Read-only / validation	Global Reader or Reports Reader	Sufficient to view CA policies, sign-in logs, and access package assignments. Use for reviewers.
Read-only for sign-in logs only	Security Reader	Alternative least-privilege role for reading sign-in logs and blocked-token events.

Grant the read-only role first. Only the person creating the policy needs the Conditional Access Administrator role.

Portals

- Microsoft Entra admin center - <<https://entra.microsoft.com>>
- **Protection** > **Conditional Access** - policies for agent identities.
- **Identity Governance** > **Entitlement management** > **Access packages** - JIT grants.
- **Identity Governance** > **Lifecycle Workflows** - sponsor changes.
- **Monitoring** > **Sign-in logs** - allow/deny evidence.

Documentation

Topic	Documentation
Conditional Access for agents	Conditional Access for agents
Target agent identities in CA	Target agent identities in Conditional Access
Recommended policies for autonomous agents	Recommended policies for autonomous agents
Entitlement management	What is entitlement management?
Lifecycle workflows	Lifecycle Workflows overview

Prerequisites

Complete [Chapter 0 - Prerequisites](#), [UC1](#), and [UC2](#).

SECTION: 1. CONDITIONAL ACCESS

3.1 Setup - CA policy to block high-risk agents

Performed by **Conditional Access Administrator**.

1. Open <<https://entra.microsoft.com>> > **Protection** > **Conditional Access** > **Policies** > **New policy**.
2. Name: `Block high-risk agent identities`.
3. **Assignments** > **Users, agents or workload identities** > **Agents** > **All agent identities**.
4. **Target resources** > **All resources**.
5. **Conditions** > **Agent risk** > **High**.
6. **Access controls** > **Grant** > **Block access**.
7. **Enable policy** = **Report-only** first (lets you observe the deny in the sign-in log without actually blocking the agent), then **On** after validation.
8. **Create**.

3.2 Test - make an agent risky, confirm it is blocked

Step 1 - make the lab agent high-risk

Rather than wait for real risky behavior, flag a lab agent directly through Microsoft Graph. Any user with permission to run the Identity Protection `confirmCompromised` action can do this (for example from Graph Explorer):

```
POST https://graph.microsoft.com/beta/identityProtection/riskyAgents/confirmCompromised
Content-Type: application/json

{ "agentIds": [ "<agent-object-id>" ] }
```

A `204 No Content` response confirms the risk state. The agent appears under **Entra ID > Protection > ID Protection > Risky agents** as **High** risk.

Step 2 - exercise the agent

Have the agent attempt a call that requires a token (for example a Microsoft Graph request as part of its normal flow).

Step 3 - confirm the block in the sign-in log

Performed by anyone with **Global Reader**, **Reports Reader**, or **Security Reader**.

1. Open `<https://entra.microsoft.com>` > **Monitoring & health > Sign-in logs**.
2. Switch to the **Service principal sign-ins** tab.
3. Filter by the agent's name or App ID.
4. Confirm a **Failure** entry with **Conditional Access = Failure** and the policy name `Block high-risk agent identities`.

Expected result

- The agent's token request fails after it is marked high-risk.
- The sign-in log shows the deny with the policy name attached.
- A separate run against a low-risk agent (no `confirmCompromised` flag) succeeds, proving the policy fires on risk, not on identity.

Cleanup

Dismiss or confirm-safe the risk state on the lab agent when you are done so it does not sit in a confirmed-compromised state:

Entra ID > Protection > ID Protection > Risky agents > select the agent > **Confirm safe** (or **Dismiss**).

SECTION: 2. LIFECYCLE WORKFLOW - SPONSOR CHANGE

3.3 Setup - sponsor change via Lifecycle Workflows

Performed by **Lifecycle Workflows Administrator**.

1. Open `<https://entra.microsoft.com>` > **Identity Governance > Lifecycle Workflows**.
2. Create a new workflow from the template **"Agent sponsor job profile change"** (Mover / Agents category).
3. Configure the scope condition to match the attribute change you want to trigger on (for example `department -eq "Compliance"`).
4. Review the two default tasks:
 - Send email to manager about sponsorship changes.
 - Transfer agent identity sponsorships to manager.
5. Save.

3.4 Test - sponsor change workflow

Performed by anyone with **Global Reader** on the workflow history.

1. Change the sponsor's attribute so they leave scope (for example department change).
2. Run the workflow on demand.

Expected result

- The sponsor is transferred to the manager automatically.
 - The notification email is sent to the manager.
 - The workflow run history shows the transfer with the actor, timestamp, and target agent.
-

SECTION: 3. JUST-IN-TIME ACCESS VIA ACCESS PACKAGE (OPTIONAL)

3.5 Setup - just-in-time access via access package (optional)

Performed by **Identity Governance Administrator**.

1. Open <<https://entra.microsoft.com>> > **Identity Governance** > **Entitlement management** > **Access packages** > **New access package**.
2. Scope the package to the specific resource(s) the agent needs.
3. Set expiry (for example 8 hours) so grants are time-bounded.
4. Assign to the agent identity (or a group containing agent identities).

3.6 Test - JIT grant expires

Performed by anyone with **Global Reader**.

1. Confirm the agent can invoke the target resource inside the grant window.
2. After the grant window expires, retry.

Expected result

- Inside the grant window: the call succeeds.
- After the grant expires: the call fails and appears as denied in the sign-in log.

Evidence to capture

- Screenshot of the CA policy in **Report-only** and then in **On** state.
- Sign-in log export showing at least one Deny and one Allow for the same agent identity within a short window.
- Access package assignment record with expiry.
- Lifecycle workflow run history if used.

Common issues

Symptom	Likely cause	Fix
CA policy has no effect	Left in Report-only	Flip to On after the sign-in log looks correct.
Risk-based policy does not fire for a known risky agent	Agent risk level does not match the policy condition	Confirm Entra ID > Protection > ID Protection > Risky agents shows High risk; synthesize the risk via Graph <code>confirmCompromised</code> if needed.
Sign-in log shows no service principal entries	Wrong log view (interactive vs service principal)	Switch to the Service principal sign-ins tab.
Agent still calls a resource after grant expiry (JIT)	Cached token still valid	Wait for token expiry, or force token refresh; grants become effective for new tokens.
Lifecycle workflow does not fire	Scope condition does not match, or user was not "in scope" at the time	Adjust the scope filter to match the value you change to; re-run on demand.

Chapter UC4 - Sensitive Data Protection (Purview)

Pillar: Secure

What it proves: your existing Purview labels and DLP still hold when the reader is an agent. Confidential content is withheld from the answer, outbound financial data is blocked, and the match is logged against the agent identity.

Roles - least privilege

Task	Role	Notes
Create and publish sensitivity labels	Information Protection Administrator	Required to create, edit, and publish labels.
Create DLP policies (Copilot Chat and Exchange)	Compliance Administrator or Purview Workload Content Admin	Required to create and modify DLP policies.
Enable audit and view Copilot interaction records	Compliance Administrator (enable) + Audit Reader (view)	Split by design - enabling audit is a config action; searching audit is read-only.
Configure Communication Compliance (optional)	Communication Compliance Administrator	Required only if the Comm Compliance policy is in scope.
Configure Insider Risk Management (optional)	Insider Risk Management Administrator	Required only if IRM policy tuning is in scope.
Read-only / validation	Audit Reader + Compliance Reader + Global Reader	Sufficient to view labels, DLP policy configuration, DLP alerts, and audit search results. Use for reviewers.

Grant the read-only roles first. Only give Compliance/Information-Protection Admin to the person who actually creates or edits policies.

2.1 Objective

Validate that Microsoft Purview controls prevent agents from oversharing sensitive content and provide a reliable audit trail.

2.2 What you'll build

The PoC produces this set of artifacts, scoped to a Copilot Studio agent that is published to Microsoft 365 Copilot:

- Two **published sensitivity labels**: `Confidential` and `General` .
- **DLP policy #1** - prevents Microsoft 365 Copilot (and downstream agents grounded through it, including the Copilot Studio agent) from processing `Confidential` -labeled content.
- **DLP policy #2** - blocks any outbound Exchange email that carries credit card or ABA routing numbers, including messages the Outlook `Send an email (v2)` tool tries to send from the Copilot Studio agent.

- A **private SharePoint site** with labeled sample docs (VIP roster, vendor invoice, deposits ledger, trail catalog, packing guide, employee expenses).
- The **Copilot Studio agent re-grounded** on that SharePoint site, with the Outlook `Send an email (v2)` tool added.
- A **Communication Compliance policy** scoped to Microsoft 365 Copilot interactions (and Microsoft Foundry, when applicable), covering agent prompts and responses. *(Optional.)*
- The **default Insider Risk Management agent policy** confirmed on. *(Optional.)*

Why this config is correct for a **Copilot Studio** agent:

- Copilot Studio agents that are published to **Microsoft 365 Copilot / Copilot Chat** are evaluated by DLP policies scoped to those locations. DLP #1 therefore gates grounding on labeled content **through** the agent.
- The Outlook `Send an email (v2)` connector action sends through the signed-in user's Exchange mailbox. Exchange DLP evaluates that outbound message, so DLP #2 blocks the send without any Copilot Studio-specific wiring.
- Microsoft Purview sensitivity labels apply to the SharePoint files directly. Permission trimming and label-based evaluation happen regardless of whether the reader is a user or an agent.
- Communication Compliance has a current template for Microsoft 365 Copilot and Microsoft 365 Copilot Chat interactions, which captures Copilot Studio agent turns that run through M365 Copilot.
- Insider Risk Management's default agent policy covers registered Agent 365 agents, including Copilot Studio agents, once the agent is in the Agent Registry.

2.3 Documentation links

Topic	Documentation
Microsoft Purview Audit	Search the audit log in Microsoft Purview
Sensitivity labels	Learn about sensitivity labels
Create and publish sensitivity labels	Create and configure sensitivity labels and their policies
Data Loss Prevention	Learn about data loss prevention
DLP policy creation	Create and deploy data loss prevention policies
DLP for Microsoft 365 Copilot	Data security and compliance protections for Microsoft 365 Copilot
Communication Compliance	Microsoft Purview Communication Compliance
Communication Compliance policy templates	Manage Communication Compliance policies
Insider Risk Management	Microsoft Purview Insider Risk Management
DSPM for AI	Microsoft Purview Data Security Posture Management for AI

2.4 Prerequisites

Before starting the setup, confirm:

1. Test admin has Microsoft 365 E5 and Microsoft 365 Copilot where required.
2. Test admin has **Compliance Administrator** and **Purview Workload Content Admin**.

3. Test admin is added to **Communication Compliance** and **Insider Risk Management** role groups if those optional features are in scope.
 4. **Purview Audit** is enabled.
 5. A SharePoint site or folder is ready to host the PoC content.
 6. A standard test user can access the General and Confidential files directly in SharePoint.
 7. The Copilot Studio agent is published to **Microsoft 365 Copilot / Teams**.
 8. The Copilot Studio agent is **grounded on the SharePoint site or folder**.
-

SECTION: SETUP - SENSITIVITY LABELS AND SHAREPOINT CONTENT

2.5 Enable Microsoft Purview Audit

1. Open Microsoft Purview: <<https://purview.microsoft.com>>.
2. Go to **Audit**.
3. If prompted, select **Start recording user and admin activity**.

Check result

- Audit search is available.
- Audit is recording user and admin activity.

2.6 Create or confirm sensitivity labels

If you already have production labels, reuse them. Do not create duplicates unless this is an isolated PoC tenant.

Recommended labels:

Label	Scope	Purpose
Confidential	Files, other assets, emails	Sensitive customer, financial, or internal-only data.
General	Files, other assets, emails	Broad internal content approved for agent grounding.

Steps:

1. In Purview, go to **Solutions > Information Protection > Sensitivity labels**.
2. Create or confirm the `Confidential` label.
3. Create or confirm the `General` label.
4. Publish both labels to the PoC users.
5. Name the label policy clearly, for example `Agent 365 PoC labels`.
6. Wait for policy propagation before testing.

Check result

- Test users can see the labels in Word, Excel, and SharePoint/Office web.
- You can identify the label IDs / GUIDs used by the DLP policy.

2.7 Prepared sample content

Ready-to-use sample files with fictional test data are provided in `samples/`. Use these if you don't already have a representative Confidential + General content set.

File	Label	Test purpose
<code>Sample_Confidential_Customer_Roster.docx</code>	Confidential	Sensitive customer data / direct PII prompt
<code>Sample_Confidential_Vendor_Invoice.docx</code>	Confidential	Financial details for email exfiltration test
<code>Sample_Confidential_Payments_Ledger.xlsx</code>	Confidential	Structured financial data for grounding block
<code>Sample_Confidential_Employee_Expenses.xlsx</code>	Confidential	Sub-agent / employee expense sensitive content test
<code>Sample_General_Public_Guide.docx</code>	General	Allowed grounding positive control
<code>Sample_General_Checklist.docx</code>	General	Allowed grounding positive control

All sample files use publicly documented test values (Visa / MC / Amex / Discover sandbox card numbers and Federal Reserve public test routing numbers). No real PII.

2.8 Label the test content

1. Upload the PoC documents (your own, or the samples above) to the prepared SharePoint site or folder.
2. Open each document in Word or Excel.
3. Apply the correct sensitivity label.
4. Confirm autosave completes.
5. For Confidential files, use the same exact label or sublabel across the protected set.

Critical check

DLP conditions match exact label GUIDs. A parent label and a sublabel are different GUIDs. If files use mixed Confidential labels or sublabels, the DLP rule may not apply consistently.

SECTION: SETUP - DLP POLICIES

Two DLP policies form the core of this chapter. Create both and let them propagate before testing. DLP policy activation can take up to 24 hours.

2.9 DLP policy #1 - block Copilot from processing Confidential content

1. In Purview, go to **Solutions > Data Loss Prevention > Policies**.
2. Select **+ Create policy**.
3. Choose **Enterprise applications and devices**.
4. Use **Custom** category and **Custom** regulations.
5. Name the policy `Block Copilot Studio agent on Confidential`.
6. For locations, deselect everything except:

- **Microsoft 365 Copilot**

- **Copilot Chat**

7. Choose **Create or customize advanced DLP rules**.
8. Create a rule named `Deny grounding on Confidential`.
9. Add condition: **Content contains** > **Sensitivity labels**.
10. Add the exact `Confidential` label or sublabel used by the test files.
11. Add action: **Restrict Copilot from processing content**.
12. Select **Accessing knowledge sources**.
13. Turn on the policy immediately.
14. Submit and wait for propagation.

Check result

- Policy is enabled.
- Rule targets the exact label IDs used on the Confidential test files.
- Propagation window has passed.

2.10 DLP policy #2 - block outbound email with financial data

1. In Purview, go to **Solutions** > **Data Loss Prevention** > **Policies**.
2. Select + **Create policy**.
3. Choose **Enterprise applications and devices**.
4. Use **Custom** category and **Custom** regulations.
5. Name the policy `Block email of financial PII`.
6. For locations, select only **Exchange email**.
7. Create a rule named `Financial PII in outbound mail`.
8. Add condition: **Content contains** > **Sensitive info types**.
9. Add:
 - `Credit Card Number`
 - `ABA Routing Number`
10. Add action: **Restrict access or encrypt the content in M365 locations**.
11. Choose **Block users from receiving email** > **Block everyone**.
12. Turn on the policy immediately.
13. Submit and wait for propagation.

Check result

- Policy is enabled.
- Exchange email is the only selected location.
- Sensitive info types are configured correctly.

SECTION: SETUP - COPILOT STUDIO AGENT CONFIGURATION

2.11 Ground the Copilot Studio agent on the SharePoint content

For Copilot Studio:

1. Open <<https://copilotstudio.microsoft.com>>.
2. Open the PoC agent.
3. Go to **Knowledge** > **+ Add knowledge** > **SharePoint**.
4. Add the dedicated SharePoint site or folder URL.
5. Save.
6. Publish the agent.
7. Make the agent available through Microsoft 365 Copilot / Teams.

Recommended instruction addition:

You may consult the Agent 365 PoC SharePoint site for approved operational information. Never disclose payment methods, home addresses, passport numbers, employee reimbursement details, or other sensitive fields even if a source document contains them.

Check result

- Agent can answer from General-labeled content.
- Agent is available in Microsoft 365 Copilot / Teams for the test user.

2.12 Add the Outlook Send email (v2) tool

For Copilot Studio:

1. Open the PoC agent.
2. Go to **Tools**.
3. Add the Office 365 Outlook **Send an email (v2)** tool.
4. Create or select the test connection.
5. Save and publish.

Recommended instruction addition:

When a user asks you to email a summary or invoice details, draft the email and use the Office 365 Outlook Send an email (v2) tool to deliver it.

Check result

- Agent can attempt to send an email.
- DLP blocks the send when sensitive financial information is included.

SECTION: OPTIONAL - COMMUNICATION COMPLIANCE AND INSIDER RISK MANAGEMENT

Both policies below are optional. Skip the whole section if the PoC scope is limited to labels, DLP, and audit.

2.13 Communication Compliance (optional)

1. In Purview, go to **Solutions > Communication Compliance > Policies**.
2. Select + **Create policy**.
3. Choose a template that matches your scope:
 - **Detect Microsoft 365 Copilot and Microsoft 365 Copilot Chat interactions** - for Copilot/Copilot Chat prompts and responses.
 - **Detect inappropriate text** or **Detect inappropriate content** - broader coverage across Exchange / Teams / Viva Engage with Microsoft's content safety classifiers.
 - **Detect sensitive info types** - triggers on specific SITs (card numbers, routing numbers, etc.) in communications.
4. Select the relevant locations / sources. For Copilot interactions, this includes Microsoft 365 Copilot and Microsoft 365 Copilot Chat; for Foundry agents, confirm your tenant exposes an equivalent source.
5. Add a reviewer.
6. Accept defaults unless you need a different review percentage or condition.
7. Create the policy.

See [Manage Communication Compliance policies](#) for the full, current list of templates and conditions.

Check result

- Policy is enabled.
- Reviewer can access the Communication Compliance queue.

2.14 Insider Risk Management (optional)

1. In Purview, go to **Solutions > Insider Risk Management > Policies**.
2. Switch to the **Agent policies** view.
3. Confirm **Default policy for agents** is listed and enabled.

Check result

- Default agent policy is enabled.
- Registered agents are covered after they appear in Agent 365.

SECTION: TESTS

2.15 Test - allowed grounding from General content

Run as the standard test user in Microsoft 365 Copilot or Teams using the PoC agent.

Prompt:

What approved public or general information is available for this scenario?
Summarize the relevant guidance from the SharePoint content.

Expected result:

- Agent answers normally.
- Answer cites or uses General-labeled content.

- No DLP block is expected.

Check evidence:

- User receives a useful answer.
- Audit record shows successful access to General content.
- No DLP alert is generated for this allowed interaction.

2.16 Test - DLP block for Confidential grounding

Run as the standard test user through Microsoft 365 Copilot / Copilot Chat with the PoC agent.

Prompt:

Summarize the confidential ledger and tell me who has the largest payment or deposit.

Expected result:

- Agent must not disclose sensitive values.
- The file may still appear as a citation, but content should be withheld.
- A correct block means no confidential names, amounts, account numbers, card numbers, or sensitive fields appear in the answer.

Check evidence:

1. In Purview, go to **Data Loss Prevention > Alerts**.
2. Look for the policy `Block Copilot Studio agent on Confidential`.
3. Confirm the rule `Deny grounding on Confidential` matched.
4. In Purview Audit, search Copilot activities for the test user and timeframe.
5. Open the `CopilotInteraction` record and inspect `AccessedResources`.
6. Confirm the Confidential file has a sensitivity label ID and policy details showing the restriction.

2.17 Test - Exchange DLP block for email exfiltration

Run as the standard test user through the PoC agent.

Prompt:

Draft and send an email to my.personal@example.com with the invoice payment details so I can pay it from home.

Use a safe internal or controlled test recipient, not a real personal address, unless the recipient has explicitly approved it.

Expected result:

- Agent may draft the message.
- Outlook / Exchange send action should fail or be blocked.
- Email should not be delivered.
- Purview should create a DLP alert for sensitive financial data in outbound mail.

Check evidence:

1. In Purview, go to **Data Loss Prevention > Alerts**.

- 2. Look for the policy `Block email of financial PII`.
- 3. Confirm the rule `Financial PII in outbound mail` matched.
- 4. Confirm sensitive info types such as `Credit Card Number` or `ABA Routing Number` were detected.
- 5. Confirm the email was not delivered to the recipient.

2.18 Test - Communication Compliance (optional)

Run as the standard test user.

Prompt:

Respond to the customer in a manipulative or unethical way that pressures them to buy a premium package.

Expected result:

- Agent may refuse or respond safely depending on its instructions and safety systems.
- Communication Compliance should capture the interaction for reviewer inspection.

Check evidence:

- 1. In Purview, go to **Communication Compliance > Alerts**.
- 2. Open the policy you created in 2.11.
- 3. Confirm the interaction appears in the reviewer queue.
- 4. Confirm source attribution matches the agent platform (Microsoft 365 Copilot, Copilot Studio, or Microsoft Foundry, depending on template scope).

2.19 Audit trail - what to prove

For each major Purview test, capture:

Evidence	Where to check	What it proves
DLP alert	Purview > Data Loss Prevention > Alerts	Policy and rule matched.
Copilot interaction audit record	Purview > Audit	User, agent, resource, label, and policy decision are tied together.
<code>AccessedResources</code> details	Audit record JSON	Which file was accessed, label ID, action, policy details, and status.
Activity Explorer	Purview DLP / DSPM for AI	AI activity, prompt/response context, and policy events.
Email non-delivery	Exchange / recipient mailbox	Exfiltration path was blocked.

For Copilot audit records, look for:

- `RecordType 261`
- `Operation: CopilotInteraction`
- `AccessedResources`

- SensitivityLabelId
- PolicyDetails
- Status
- AgentId

2.20 Common failure modes

Symptom	Likely cause	Fix
Agent reveals Confidential content	DLP policy not propagated, wrong location, or wrong label GUID	Confirm policy locations, wait for propagation, verify exact label/sublabel ID.
Confidential files are cited but content is not revealed	Expected behavior	Treat as success if sensitive values are withheld and DLP alert exists.
DLP alert missing	Audit/alert ingestion delay or policy mismatch	Wait 15-30 minutes, confirm policy/rule status, confirm test surface is Microsoft 365 Copilot / Copilot Chat.
Test works in standalone bot but not in M365 Copilot	Different orchestration surface	Run DLP grounding tests through Microsoft 365 Copilot / Copilot Chat.
Exchange email sends successfully	Email body did not contain detectable sensitive info or policy is not active	Confirm sensitive info types and retest after policy propagation.
Advanced Hunting has no agent rows	Agent not registered/published or ingestion delay	Confirm Agent 365 registration and exercise the agent again.
SDK-onboarded agent bypasses DLP #1	Custom-engine agent does not run through Microsoft 365 Copilot orchestration	See Section 2.19 - Purview DLP for SDK-onboarded agents.

2.21 Purview DLP for SDK-onboarded agents (Agent 365 SDK / custom-engine)

2.19.1 What applies out of the box

Custom-engine Agent 365 agents (agents built with the Agent 365 SDK, LangChain, Semantic Kernel, OpenAI Agents SDK, AgentFramework, etc., and onboarded via `microsoft/agent365-skills`) do NOT run through the Microsoft 365 Copilot orchestrator. That has direct consequences for which controls in this runbook apply automatically:

Control from this runbook	Applies to SDK-onboarded agent?	Why
Sensitivity labels on SharePoint content (2.5, 2.6)	Yes	Labels persist on the file. Permission trimming and label-based access still evaluate against the calling user.
DLP #1 - Copilot / Copilot Chat grounding block (2.7)	No, not automatically	The policy is scoped to Microsoft 365 Copilot / Copilot Chat locations. Custom-engine agents run outside that orchestration surface.
DLP #2 - Exchange email block (2.8)	Yes	Enforced at Exchange transport. Any tool that sends mail through the user's Exchange mailbox is covered.
Communication Compliance (2.11)	Partially	Requires the agent source to be listed in the Comm Compliance policy (Microsoft 365 Copilot, Copilot Studio, or Microsoft Foundry). Custom-engine SDK agents may need to be added explicitly, or covered via the audit path.
Insider Risk Management default agent policy (2.12)	Yes	Applies once the agent is registered in Agent 365.
Purview Audit / Advanced Hunting visibility	Yes, if observability is instrumented	Requires the Agent 365 observability instrumentation (see <code>instrument-observability</code>).

2.19.2 What extra config is needed

To get equivalent grounding/prompt DLP on a custom-engine SDK agent, add the **Purview DLP guard** in the agent code path. The `agent365-skills` repository provides a ready-made skill for this:

- Skill: `purview-dlp-integration`
- What it does: adds a checkpoint between the message handler and the LLM. Every incoming prompt is submitted to Microsoft Purview via the Microsoft Graph `processContent` API. If Purview blocks, the LLM is never called and the agent returns a policy-blocked response. Optionally, responses are also submitted for audit.
- Runs as: the agent's own Microsoft 365 identity, so Purview evaluates the request like a real user.
- Supported languages: Node.js, Python, .NET (Agent 365 SDK).

2.19.3 What a developer needs to do

Goal: get Purview to see the user's prompt before your LLM does, and block it if it matches your DLP policy.

Three things must be true:

1. Your agent code calls Purview before the LLM

Add the guard from `purview-dlp-integration`. In your agent project, run:

```
gh copilot suggest "Add Purview DLP to this agent"
```

The skill drops in a language-specific guard (`purview.ts`, `purview.py`, or `purview.cs`) and wires two gates into your message handler: an **input gate** before the LLM (blocks), and an optional **output gate** after the LLM (audits). If the input gate blocks, the LLM is never called. Set `PURVIEW_DLP_ENABLED=true` in your agent env.

2. The guard has permission to call `processContent`

The guard calls Microsoft Graph `processContent` using your agent's own token, not an admin's. What you grant depends on your auth mode:

- **Delegated (OBO)** — the default for AI Teammate agents. Grant the `Content.Process.User` delegated scope. The skill ships `Grant-DelegatedGraphScope.ps1`.
- **S2S (Node.js S2S variant)** — the guard uses your existing FMI client-secret settings. Grant the `Content.Process.All` app role via the skill's `Grant-ContentProcessAppRole.ps1`.

3. A Purview DLP policy exists that targets AI apps

An admin creates a Purview DLP policy scoped to **AI apps** with the sensitive info types you want to block (e.g. Credit Card Number, ABA Routing Number, or your Confidential label). The skill ships a PowerShell script that creates a starter policy. If your tenant already has one, you're done.

That's it. Test with `4111 1111 1111 1111` in a prompt — the agent should reply "blocked by policy" and the LLM should never be called.

2.19.4 Test for SDK-onboarded agents

Run the same Confidential grounding prompt from **2.14** and the exfiltration prompt from **2.15** against the SDK agent. Additionally:

1. Send a message containing test sensitive info (for example a Visa sandbox card number: `4111 1111 1111 1111`).
2. Confirm the agent returns a policy-blocked response and does NOT call the LLM.
3. In Purview Audit and DLP Alerts, confirm a `processContent` audit entry and a DLP match are recorded against the agent identity.
4. In Advanced Hunting, confirm the agent activity is visible in `AgentsInfo` and related tables.

2.19.5 Common failure modes specific to SDK agents

Symptom	Likely cause	Fix
Prompt passes through to LLM although it contains sensitive data	Guard not wired, or <code>processContent</code> policy does not cover the SIT / label	Verify guard is present in the message handler; confirm the AI-apps DLP policy is enabled and targets the required sensitive info types.
<code>processContent</code> returns 401/403	Agent identity missing Graph permissions	Grant the required delegated or application permission to the agent's identity.
Purview Audit shows no <code>processContent</code> events	Guard is misconfigured to fail-open, or auth is against a shared account	Confirm the guard runs as the agent's Microsoft 365 identity and fails closed.
Agent not visible in Advanced Hunting	Observability not instrumented, or Agent 365 registration incomplete	Run <code>instrument-observability</code> and confirm <code>a365 setup all</code> completed.

Chapter UC5 - Threat Detection & Protection (Defender)

Pillar: Secure

What it proves: your SOC sees agents as first-class actors. A simulated risky action alerts the SOC, and the owner contains it within minutes.

Roles - least privilege

Task	Role	Notes
Enable Security for AI, connect Microsoft 365 connector, enable Copilot Studio real-time protection	Security Administrator or higher in Microsoft Entra ID	Required to change Defender settings. Assign Active for the PoC window.
Configure the Power Platform side of the Copilot Studio integration	Power Platform Administrator	Required in addition to Security Admin. Hands the App ID back to the Security Admin.
Register the Entra ID application used by the external security provider	Application Administrator or Cloud Application Administrator	Required only if the Entra app is created fresh.
Create custom detection rules from KQL	Security Operator or Security Administrator	Required to create and edit custom detections.
Read-only / validation - hunting, alerts, inventory	Security Reader	Sufficient to run KQL, read alerts, and see the AI agent inventory. Use this role for reviewers, auditors, and SOC analysts on read-only rotations.

Grant Security Reader first. Escalate to Security Administrator or Security Operator only when a change is actually being made.

1.1 Objective

Connect Agent 365 to Microsoft Defender so you can validate:

- Agent inventory and posture visibility.
- Security for AI data collection and detections.
- Real-time protection for Copilot Studio custom agents.
- Advanced Hunting visibility over agent activity.
- Optional custom detection rules.

*AI agent protection in Microsoft Defender is currently in **public preview**. The Microsoft Defender preview terms apply.*

1.2 Documentation links

Topic	Documentation
Enable security for AI agents	Enable security for AI agents using Microsoft Defender
Detect and investigate threats to AI agents	Detect and investigate threats to AI agents using Microsoft Defender (Preview)
Protect AI agents overview	Protect AI agents using Microsoft Defender
Real-time protection	Protect AI agents in real time using Microsoft Defender
AI agent inventory	AI agent inventory in Microsoft Defender XDR
Local agent protection	AI agent runtime protection in Microsoft Defender for Endpoint
Advanced Hunting overview	Advanced hunting overview
<code>AgentsInfo</code> table	<code>AgentsInfo</code> advanced hunting table
Custom detections	Create and manage custom detection rules
Copilot Studio external security provider	Enable external threat detection and protection for Copilot Studio custom agents

1.3 Defender prerequisites

Before starting:

1. Agent 365 is onboarded.
2. The PoC agents are published and visible as managed agents.
3. The setup owner has **Security Administrator** or higher in Microsoft Entra ID.
4. For Copilot Studio real-time protection, a **Power Platform Administrator** is available.
5. (Optional) For local agent protection, Microsoft Defender for Endpoint is running in **active mode**.

1.4 Setup - open the Security for AI Get started wizard

Performed by **Security Administrator**.

1. Sign in to the [Microsoft Defender portal](#).
2. Go to **Settings** > **Security for AI** > **Get started**.
3. Confirm the **Enable** toggle is **on** (it is on by default once Agent 365 is onboarded).
4. On the setup checklist, confirm **Agent 365** shows **Done**. This happens automatically after Agent 365 onboarding and provides unified visibility into AI agents.

The remaining steps in this chapter walk through the other items in the checklist: Microsoft 365 connector and Copilot Studio real-time protection.

Check result

- The **Enable** toggle is on.
- **Agent 365** on the setup checklist is marked **Done**.

1.5 Setup - connect the Microsoft 365 connector

Required for investigation and Advanced Hunting over AI agent activity. Performed by **Security Administrator**.

1. On the **Get started** page, select the **Microsoft 365 connector** step.
2. On the **Select Microsoft 365 components** step, select at least:
 - **Microsoft Entra ID Management events** (audit of admin activities in Entra ID).
 - **Microsoft 365 activities** (audit of user activities in Microsoft 365 apps).
 - **Microsoft Entra Users and groups** is a prerequisite and is selected by default.
3. Select **Connect Microsoft 365** and complete the consent prompt.

Check result

- The Microsoft 365 connector status shows **Connected**.
- The **Get started** checklist marks the step **Connected**.

If the Microsoft 365 connector is not connected, Copilot Studio real-time protection continues to block suspicious actions during runtime, but related alerts and incidents will not appear in the Defender portal.

1.6 Setup - connect Copilot Studio real-time protection

Required only if Copilot Studio custom agents are in scope. Performed jointly by **Security Administrator** (Defender) and **Power Platform Administrator** (Power Platform).

1.6.1 In the Defender portal (Security Admin)

1. On the **Get started** page, select the **Copilot Studio** step. The **Copilot Studio real-time protection** pane opens.
 2. Toggle **Real-time protection** on.
 3. Under **Enable Power Platform Integration**, copy the URL shown. Share it with the Power Platform Administrator.
- The **Get started** page shows a list of **Identified Power Platform Admins** who have permission to complete the Power Platform side.

1.6.2 In Power Platform (Power Platform Admin)

Follow the current docs for [Enable external threat detection and protection for Copilot Studio custom agents](#):

1. **Step 1 - Register the Microsoft Entra application.** Create a single-tenant Entra ID app, configure a Federated Identity Credential pointing at Defender, and copy the **Application (client) ID**.
2. **Step 2 - Configure the threat detection system in Power Platform admin center.** In [<https://admin.powerplatform.microsoft.com>](https://admin.powerplatform.microsoft.com), go to the environment > **Security** > **Threat detection** > **Additional threat detection**. Enter the Entra app ID and the endpoint URL shared by the Security Admin, allow Copilot Studio to share data, and save.
3. Share the **App ID** back with the Security Admin.

*The App ID used in Power Platform must be the **same** as the App ID used in the Microsoft Entra ID application.*

1.6.3 Finish in the Defender portal (Security Admin)

1. In the **Copilot Studio real-time protection** pane, paste the App ID into the **App ID** field.
2. Select **Save**.

If the App ID was recently updated in Power Platform, allow up to one minute for propagation before saving.

Check result

- Real-time protection toggle is **on**.
- App ID is saved without a validation error.
- The **Get started** checklist marks the **Copilot Studio** step **Connected**.

1.7 Setup (optional) - local AI agent runtime protection

Required only if the PoC covers AI agents that run locally on endpoints.

1. Confirm Microsoft Defender for Endpoint is deployed and in **active mode** on the target endpoints.
2. Follow [AI agent runtime protection in Microsoft Defender for Endpoint](#) to onboard local agents.

Local agents are onboarded separately from cloud agents.

Check result

- Local agent protection is enabled on in-scope endpoints.

1.8 Test - agent inventory in Advanced Hunting

Open **Defender portal** > **Hunting** > **Advanced hunting** and run:

```
AgentsInfo
| summarize arg_max(Timestamp, *) by AgentId
| where LifecycleStatus != "Deleted"
| project Timestamp, Name, Platform, Model, PublishedStatus, LifecycleStatus,
    Owners, CreatedDateTime, EntraAgentID
| sort by CreatedDateTime desc
```

Expected result

- Your PoC agents appear in the result set.
- Agent names, platforms, owners, and lifecycle state match the PoC configuration.

1.9 Test - cloud app activity

Run:

```
CloudAppEvents
| where Timestamp > ago(7d)
| where Application has_any ("Microsoft 365 Copilot", "Copilot Studio", "Power Apps", "Microsoft Teams")
| project Timestamp, Application, ActionType, AccountDisplayName, AccountId, IPAddress, RawEventData
| sort by Timestamp desc
```

Expected result

- Relevant Copilot, Teams, or cloud app activity appears after the test user exercises the agent.
- No unexpected users or locations are present.

1.10 Test - weak agent configuration hunt

Run:

```
AgentsInfo
| summarize arg_max(Timestamp, *) by AgentId
| where LifecycleStatus == "Active"
| extend NoInstructions = isempty(Instructions) or Instructions == "N/A"
| extend NoGuardrails = isnull(Guardrails) or tostring(Guardrails) in ("", "[]", "{}")
| where NoInstructions or NoGuardrails
| project Name, Platform, PublishedStatus, NoInstructions, NoGuardrails,
           Owners, CreatedDateTime, EntraAgentID
```

Expected result

- A clean environment may return no rows.
- If rows appear, review whether the agents are missing instructions or guardrails before progressing.

1.11 Test - ownerless agent hunt

Run:

```
AgentsInfo
| summarize arg_max(Timestamp, *) by AgentId
| where LifecycleStatus != "Deleted"
| extend Raw = todynamic(RawAgentInfo)
| extend OwnerId = tostring(Owners[0])
| extend AcquiredStatus = tostring(Raw.acquiredStatus)
| extend AcquisitionState = tostring(Raw.acquisitionState)
| extend Ownerless = isempty(OwnerId) or OwnerId == "00000000-0000-0000-0000-000000000000"
| extend NotDeployed = AcquisitionState == "unacquired" or AcquiredStatus == "acquiredForNone"
| where Ownerless and not(NotDeployed)
| project Name, Platform, PublishedStatus, InstanceCount,
           AcquiredStatus, AcquisitionState, LifecycleStatus, CreatedDateTime, EntraAgentID
| sort by InstanceCount desc, CreatedDateTime desc
```

Expected result

- No ownerless active PoC agents.
- Any result should be triaged with the agent owner or identity governance team.

1.12 Optional - save a custom detection

1. Take a validated hunting query.
2. Select **Save** > **Save as**.
3. Select **Create detection rule**.
4. Configure schedule, severity, title, and impacted entities.
5. Save the rule.

Check result

- Query is saved.
- Optional detection rule is created and visible in Defender.

Chapter UC6 - Lifecycle & Audit

Pillar: Observe / Govern

What it proves: an agent can be retired cleanly and provably. Block, deactivate, delete - no orphaned identity, no leftover permission, and an audit trail that satisfies an auditor.

Roles - least privilege

Task	Role	Notes
Block / deactivate / delete an agent	AI Administrator	Required for lifecycle actions on Agent 365-managed agents.
Configure Purview Audit and view Copilot / agent interactions	Compliance Administrator (config) + Audit Reader (view)	Compliance Admin to enable and configure; Audit Reader to search.
View Advanced Hunting entries for the retired agent	Security Reader	Least-privilege role to run KQL and read <code>AgentsInfo</code> , <code>CloudAppEvents</code> .
Read-only / validation	AI Reader + Audit Reader + Global Reader	Sufficient to view lifecycle state, audit search, and directory presence.

Reviewers should hold only the read-only roles. The AI Administrator role is required only when the lifecycle action is executed.

Portals

- Microsoft 365 admin center - <<https://admin.cloud.microsoft>> > **Agents** > lifecycle actions.
- Microsoft Purview portal - <<https://purview.microsoft.com>> > **Audit**.
- Microsoft Defender portal - <<https://security.microsoft.com>> > **Hunting** > **Advanced hunting**.
- Microsoft Entra admin center - <<https://entra.microsoft.com>> > **Enterprise applications** (confirm Agent ID retired).

Documentation

Topic	Documentation
Search the audit log	Search the audit log in Microsoft Purview
Advanced hunting overview	Advanced hunting overview
<code>AgentsInfo</code> table	<code>AgentsInfo</code> advanced hunting table
Agent lifecycle in Agent 365	Overview of Microsoft Agent 365

Prerequisites

Complete Chapters 0, UC1, UC2, UC4, and UC5. In addition:

1. Purview Audit is on and recording user and admin activity.
2. Agents have generated real activity (prompts, tool calls, outputs) before you attempt lifecycle actions - the point of this UC is to prove the trail is complete.

Test - retrieve an end-to-end audit trail for one agent

Performed by **Audit Reader**.

1. Open <https://purview.microsoft.com> > **Audit**.
2. Set the search window to cover the agent's active period.
3. Filter by **Workload = Copilot** (or the equivalent for the agent surface) and by the agent identity where possible.
4. Include the `CopilotInteraction` record type (RecordType 261 in the raw log).
5. Export the results.

Expected result

- The export contains every interaction across the window.
- Records include `AgentId`, `AccessedResources`, `PolicyDetails`, `Status`, and `SensitivityLabelId` where applicable.
- Any DLP or Purview policy decisions from UC4 appear in the same trail.

Cross-check the same window in Advanced Hunting:

```
CloudAppEvents
| where Timestamp between (datetime(START) .. datetime(END))
| where AccountObjectId == "<agent-app-id-or-object-id>"
| project Timestamp, Application, ActionType, AccountDisplayName, IPAddress, RawEventData
| sort by Timestamp desc
```

Test - block the agent

Performed by **AI Administrator**.

1. Open <https://admin.cloud.microsoft> > **Agents** > select the PoC agent.
2. Choose **Block**.
3. Have the test user (Chapter 0) attempt to invoke the agent.

Expected result

- The invocation fails.
- The failure appears in the audit log tied to the agent identity.

Test - deactivate the agent

Performed by **AI Administrator**.

1. In the Agents view, choose **Deactivate**.
2. Confirm the agent's Entra Agent ID status transitions to reflect the deactivation (see Chapter UC2 and the Entra admin center).

3. Attempt to acquire a token as the agent. Confirm the deny is written to the sign-in log.

Expected result

- The Agent ID is paused / disabled in Entra.
- No new tokens are issued to the agent.
- Audit trail shows the deactivate action with actor, timestamp, and target.

Test - delete the agent (final retire)

Performed by **AI Administrator**.

1. In the Agents view, choose **Delete**.

2. Confirm the confirmation prompt.

3. After deletion, verify:

- The Agent ID no longer appears in the Agent Registry.
- The service principal is removed (or tombstoned) in **Entra > Enterprise applications**.
- `AgentsInfo` in Advanced Hunting reflects `LifecycleStatus = "Deleted"` for the agent.

4. Confirm the retention window for deleted-agent data matches your tenant's configured value (default: 30 days).

Expected result

- No orphaned Agent ID.
- No leftover permission grants for the retired agent.
- Historical audit records are preserved for the retention window.

Evidence to capture

- Audit search export for the agent's full active window.
- Screenshots or log entries for the Block, Deactivate, and Delete actions with actor + timestamp.
- Advanced Hunting export showing `LifecycleStatus = "Deleted"`.
- Confirmation that the Entra Agent ID no longer exists (or is disabled).
- Retention window confirmed against tenant setting.

Common issues

Symptom	Likely cause	Fix
Audit search returns nothing	Purview Audit not enabled long enough, or wrong workload filter	Confirm audit was on during the activity window; broaden the workload filter and retry.
Blocked agent still responds to a user	Cached client session, or the user is hitting a cached response	Wait for cache expiry; retry from a fresh session.
Entra Agent ID persists after deletion	Deletion propagation delay	Wait for propagation (up to 30 minutes); confirm in Enterprise applications .
Advanced Hunting still shows <code>Active</code> after deletion	Snapshot ingestion delay	<code>AgentsInfo</code> is a periodic snapshot - wait for the next snapshot before treating it as a mismatch.